

COT 5615 Math for Intelligent Systems Fall 2021 Homework #7

UFID: 96703101

Name: *Vyom Pathak*

Instructor: Professor Kejun Huang

Due Date: November 07, 2021

Problem 10.16**Covariance matrix****Solution**

(a) $\mu = \frac{A^T \cdot 1}{n}$

(b) As it is the de-meaned value, $\tilde{A} = A - 1\mu^T$. Here, $\tilde{a}_i = a_i - \mu_i 1$, thus \tilde{A} can be shown as follows:

$$\begin{aligned}
\tilde{A} &= (a_1; a_2; \dots; a_k) - (\mu_1 1; \mu_2 1; \dots; \mu_k 1) \\
&= (a_1; a_2; \dots; a_k) - 1(\mu_1; \mu_2; \dots; \mu_k) \\
&= A - 1\mu^T
\end{aligned}$$

(c) The diagonal entries \sum_{ii} are given as follows:

$$\begin{aligned}
\sum_{ii} &= \frac{(\tilde{A}^T \tilde{A})_{ii}}{N} \\
&= \frac{\tilde{a}_i \tilde{a}_i^T}{N} \\
&= \frac{\|\tilde{a}_i\|^2}{N} \\
&= \text{std}(a_i)^2 \left[\because \text{standard deviation } \text{std}(a_i) = \frac{\|\tilde{a}_i\|}{\sqrt{N}} \right]
\end{aligned}$$

Now, the off-diagonal entries of \sum_{ij} is zero where $\tilde{a}_i = 0$ and $\tilde{a}_j = 0$. For other values, it is calculated as follows:

$$\begin{aligned}
\sum_{ij} &= \frac{(\tilde{A}^T \tilde{A})_{ij}}{N} \\
&= \frac{\tilde{a}_i \tilde{a}_j^T}{N} \\
&= \rho_{ij} \text{std}(a_i) \text{std}(a_j) \left[\because \rho_{ij} = \frac{\tilde{a}_i \tilde{a}_j^T}{N \cdot \text{std}(a_i) \text{std}(a_j)} \right]
\end{aligned}$$

(d) $Z = (A - 1\mu^T) \text{diag}(\frac{1}{\text{std}(a_1)}; \frac{1}{\text{std}(a_2)}; \dots \frac{1}{\text{std}(a_k)})$. Here, $z_i = \frac{\tilde{a}_i}{\text{std}(a_i)}$ are the standardized vectors. Thus the above solution can be derived as follows:

$$\begin{aligned}
Z &= \tilde{A} \cdot \text{diag}(\frac{1}{\text{std}(a_1)}; \frac{1}{\text{std}(a_2)}; \dots \frac{1}{\text{std}(a_k)}) \\
&= (A - 1\mu^T) \cdot \text{diag}(\frac{1}{\text{std}(a_1)}; \frac{1}{\text{std}(a_2)}; \dots \frac{1}{\text{std}(a_k)})
\end{aligned}$$

Problem A14.3

Iris classification

Solution

The following code is used to find the least square solution for each class of the iris data and also the multi-class classification model is built by using the single regression model together :

```
1      # Using this imports for all the solutions.
2      using Statistics
3      using LinearAlgebra
4      using Plots
5      using MLJ
6      plotly()
7
8      include("iris_flower_data.jl")
9      include("iris_multiclass_helpers.jl")
10     # function to get and format dataset
11     function get_iris_data()
12         Random.seed!(5)
13         # Formatting dataset (must Pkg.add("RDatasets"))
14         iris = dataset("datasets", "iris")
15         data = Matrix{iris}
16         X = transpose(1.0*data[:, 1:4]);
17         perm = randperm(150);
18         X = X[:, perm];
19         y = [ones(50); 2*ones(50); 3*ones(50)];
20         y = y[perm];
21         return (X, y)
22     end
23
24     D = get_iris_data();
25     X_train = D[1][1:4,1:100]';
26     Y_train = D[2][1:100];
27
28     Theta1 = X_train\((2*(Y_train.==1)).-1);
29     Theta2 = X_train\((2*(Y_train.==2)).-1);
30     Theta3 = X_train\((2*(Y_train.==3)).-1);
31
32     X_test = D[1][1:4,101:150]';
33     Y_test = D[2][101:150];
34
35     #a
36     #Train errors
37     display(mean((X_train*Theta1.>0) .!=(Y_train.==1)));
38     display(mean((X_train*Theta2.>0) .!=(Y_train.==2)));
39     display(mean((X_train*Theta3.>0) .!=(Y_train.==3)));
40
41     #Test errors
42     display(mean((X_test*Theta1.>0) .!=(Y_test.==1)));
43     display(mean((X_test*Theta2.>0) .!=(Y_test.==2)));
44     display(mean((X_test*Theta3.>0) .!=(Y_test.==3)));
45
46     #b
47     #confusion_matrix for train and test
```

```
48 display(confusion_matrix(argmax_by_row((X_train*[Theta1 Theta2 Theta3])), Y_train));
49 display(confusion_matrix(argmax_by_row((X_test*[Theta1 Theta2 Theta3])), Y_test));
```

- (a) Class 1 error rate for train data: 0.0
Class 2 error rate for train data: 0.28
Class 3 error rate for train data: 0.36

Class 1 error rate for test data: 0.0
Class 2 error rate for test data: 0.24
Class 3 error rate for test data: 0.28

- (b) The confusion matrix for training data (upper) and testing data (lower) is shown in the figure 1:




Figure 1: Confusion Matrix for training data (upper) and testing data for iris multi-class classification

Problem A15.2

Least squares classification with regularization

Solution

```
1 include("lsq_classifier_data.jl")
2
3 #a
4 A = [ones(size(X)[2]) X'];
5
6 Theta = A \ y;
7 v = Theta[1,:];
8 beta = Theta[2:51];
9 println("Beta and v (slope) values are : ", beta, v);
10
11 y_hat = X' * beta .+ v;
12 test_y_hat = X_test' * beta .+ v;
13 println("The error rate on training data is ", (sum(((y_hat).>0) .!=(y.==1)) + sum(((y_hat).<0) .!=(y.== -1))))
14 /size(y)[1]/2);
15 println("The error rate on testing data is ", (sum(((test_y_hat).>0) .!=(y_test.==1)) + sum(((test_y_hat).<0) .!=(
16 (y_test.== -1))))/size(y_test)[1]/2);
17
```

```

18  #b
19  regularization_factor = 10.^ range(-1,4,length = 100);
20  test_values = []
21  train_values = []
22  logvalues = []
23  for i in regularization_factor
24      Theta = inv(transpose(A)*A + i*Matrix{Float64}(I,51,51)) * transpose(A) * y;
25      v = Theta[1,:];
26      beta = Theta[2:51];
27      y_hat = X' * beta .+ v;
28      test_y_hat = X_test' * beta .+ v;
29      append!(train_values,(sum(((y_hat).>0) .!=(y.==1)) + sum(((y_hat).<0) .!=(y.==-1)))/size(y)[1]/2);
30      append!(test_values,(sum(((test_y_hat).>0) .!=(y_test.==1)) + sum(((test_y_hat).<0) .!=(
31      (y_test.==-1)))/size(y_test)[1]/2);
32      append!(logvalues,log10(i));
33  end
34
35  plot(logvalues, xlabel = "Log Scaled Lambdas", ylabel = "RMS Error", train_values, label = "Train Data")
36  plot(logvalues, xlabel = "Log Scaled Lambdas", ylabel = "RMS Error", test_values, label = "Test Data")
37  println("Regularization value at minimum loss value for train data is ", regularization_factor[findmin(train_values)[2]]);

```

1. The error rate is as follows:

The error rate on training data is 0.2

The error rate on testing data is 0.28

The beta and v (slope) values follows:

(-0.013368564887362284, 0.043914619202142506, -0.029703277493729204, 0.0441447412881381, -0.07058131397505582, 0.007607031569404489, 0.14833182533722705, 0.02852393154963966, 0.019812521586858763, -0.12070927744292764, 0.013784598389068807, 0.07897384306339864, -0.016767001698145247, 0.03105477750315689, 0.006356245285928464, -0.09111311295164341, 0.0046968705247653515, -0.03371811652421457, 0.014237165383124077, -0.1715950849684703, 0.05418932576076964, -0.06405036863460692, -0.0446363886361793, 0.060640599059827535, 0.022398446473252752, 0.01634655578292057, 0.028027646705782807, 0.046686592731404526, -0.13392652651935205, -0.09059400055804874, -0.0879007600259256, 0.010548532354429147, -0.03467602240482772, -0.06877162458039814, -0.05169326750991447, 0.13547906032991633, 0.07853780188316374, 0.11922412326678362, 0.026479714104205146, -0.053715275373795564, -0.1858378466312867, -0.06173600519777898, -0.05090661312828648, -0.0026805843235918987, 0.007823258649938056, -0.02354617634571561, 0.13831311716195221, 0.02153524018622994, -0.0036408708130813467, -0.10621012236954906) (0.3679341344425215)

2. Regularization value at minimum loss value for test data is 23.644894126454073. The figure 2

Problem A15.3

Estimating the elasticity matrix

Solution

Give the price, demand data, it can be easily transformed into a least square problem, where $\delta^d \approx \hat{E}\delta^p$ is the modeling equation. Here, first we find the delta values using the given formula. After, that we estimate elastic matrix by performing least square regression for each price value, and thus get a NXN elastic matrix. We the test this elastic matrix by using it to find change in demand for unseen data and check it's efficacy by calculating RMSE value. We perform a 80-20 split in 75 items i.e. 60 items for training and 15 items for testing. We find that, the RMSE for test and train unseen data is given as 0.208629686471379 and 0.26213391380879586. This indicates that the elastic matrix is performing well in terms of finding change in demand for new (unseen) change in prices and thus it is not showing any amount of overfitting.




Figure 2: Train and test regularization loss versus the lambda value

```

1  include("price_elasticity.jl")
2
3  delta_p = zeros(5,75)
4  delta_d = zeros(5,75)
5  for product in 1:5
6      delta_p[product,:] = (Prices[product,:].- p_nom[product]).-/p_nom[product];
7      delta_d[product,:] = (Demands[product,:].- d_nom[product]).-/d_nom[product];
8  end
9
10 delta_p_train = delta_p[1:5,1:60];
11 delta_p_test = delta_p[1:5,61:75];
12 delta_d_train = delta_d[1:5,1:60];
13 delta_d_test = delta_d[1:5,61:75];
14
15 E_cap = zeros(5,5)
16
17 for product in 1:5
18     E_cap[product,:] = (delta_d[product,:]\transpose(delta_p))
19 end
20
21 display(E_cap);
22 display(E);
23
24 delta_d_hat = E_cap * delta_p_test
25 println("The RMS Error for the training change in demand is ", rms(E_cap * delta_p_train, delta_d_train));
26 println("The RMS Error for the unseen change in demand is ", rms(delta_d_hat, delta_d_test));

```