

Homework7 solution

November 8, 2021

0.0.1 10.16

(a) $\mu = (1/n)A^T \mathbf{1}$.

This follows from $\mu_i = (\mathbf{1}^T a_i) / n = (a_i^T \mathbf{1}) / n$

(b) $\tilde{A} = A - \mathbf{1}\mu^T$.

The i th column of \tilde{A} is $\tilde{a}_i = a_i - \mu_i \mathbf{1}$, so

$$\begin{aligned}\tilde{A} &= \begin{bmatrix} a_1 & \cdots & a_k \end{bmatrix} - \begin{bmatrix} \mu_1 \mathbf{1} & \cdots & \mu_k \mathbf{1} \end{bmatrix} \\ &= \begin{bmatrix} a_1 & \cdots & a_k \end{bmatrix} - \mathbf{1} \begin{bmatrix} \mu_1 & \cdots & \mu_k \end{bmatrix} \\ &= A - \mathbf{1}\mu^T\end{aligned}$$

(c) The diagonal entries of Σ_{ii} are

$$\Sigma_{ii} = \frac{1}{N} \left(\tilde{A}^T \tilde{A} \right)_{ii} = \frac{1}{N} \tilde{a}_i \tilde{a}_i^T = \frac{1}{N} \|\tilde{a}_i\|^2 = \mathbf{std}(a_i)^2$$

since the standard deviation is defined as $\mathbf{std}(a_i) = \|\tilde{a}_i\| / \sqrt{N}$.

The off-diagonal entry Σ_{ij} is zero if $\tilde{a}_i = 0$ or $\tilde{a}_j = 0$. Otherwise,

$$\Sigma_{ij} = \frac{1}{N} \left(\tilde{A}^T \tilde{A} \right)_{ij} = \frac{1}{N} \tilde{a}_i \tilde{a}_j^T = \rho_{ij} \mathbf{std}(a_i) \mathbf{std}(a_j)$$

since the correlation coefficient is defined as

$$\rho_{ij} = \frac{1}{N} \frac{\tilde{a}_i^T \tilde{a}_j}{\mathbf{std}(a_i) \mathbf{std}(a_j)}$$

(d) $Z = (A - \mathbf{1}\mu^T) \mathbf{diag}(1/\mathbf{std}(a_1), \dots, 1/\mathbf{std}(a_k))$. The standardized vectors are $z_i = \tilde{a}_i / \mathbf{std}(a_i)$. Therefore

$$\begin{aligned}Z &= \tilde{A} \mathbf{diag}(1/\mathbf{std}(a_1), \dots, 1/\mathbf{std}(a_k)) \\ &= (A - \mathbf{1}\mu^T) \mathbf{diag}(1/\mathbf{std}(a_1), \dots, 1/\mathbf{std}(a_k)).\end{aligned}$$

0.0.2 14.3

```
[1]: using LinearAlgebra
      include("iris_flower_data.jl");
```

```

include("iris_multiclass_helpers.jl");

function least_sqaure_classifier(X, y, k)
    sz = size(X)[2]
    one = ones(sz)
    A = [X' one]
    newy = 2(y .== k) .- 1
    theta = pinv(A)*newy

    return theta, error_rate(X, y, k, theta, sz)
end

function error_rate(X, y, k, theta, sz)
    one = ones(sz)
    A = [X' one]
    y_hat = sign.(A*theta)

    newy = 2(y .== k) .- 1
    correct = sum(y_hat.==newy)
    return 1-correct/sz
end

X_train = X[:,1:100];
X_test = X[:,101:150];
y_train = y[1:100];
y_test = y[101:150];

```

(a)

```

[2]: theta1, er_train1 = least_sqaure_classifier(X_train, y_train, 1);
theta2, er_train2 = least_sqaure_classifier(X_train, y_train, 2);
theta3, er_train3 = least_sqaure_classifier(X_train, y_train, 3);
er_test1 = error_rate(X_test, y_test, 1, theta1, 50);
er_test2 = error_rate(X_test, y_test, 2, theta2, 50);
er_test3 = error_rate(X_test, y_test, 3, theta3, 50);
println("Class\t\tError rate on train\tError rate on \t
→test\t\t\t\t\t",er_train1,"\t\t\t\t", er_test1,\t
→"\tnversicolor\t",er_train2,"\t\t\t\t", er_test2,\t
→"\tnvirginica\t",er_train3,"\t\t\t\t", er_test3);

```

Class	Error rate on train	Error rate on test
setosa	0.0	0.0
versicolor	0.28	0.24
virginica	0.09999999999999998	0.0200000000000000018

(b)

```

[3]: one_train = ones(size(X_train)[2]);
one_test = ones(size(X_test)[2]);

```

```

A_train = [X_train' one_train];
A_test = [X_test' one_test];
theta = [theta1 theta2 theta3];
yy_train = argmax_by_row(A_train*theta);
yy_test = argmax_by_row(A_test*theta);
C_train = confusion_matrix(yy_train, y_train);
C_test = confusion_matrix(yy_test, y_test);

println("Confusion matrix on train set")
C_train

```

Confusion matrix on train set

```

[3]: 3×3 Matrix{Float64}:
 29.0  0.0  0.0
  0.0 24.0  5.0
  0.0 11.0 31.0

```

```

[4]: println("Confusion matrix on test set")
C_test

```

Confusion matrix on test set

```

[4]: 3×3 Matrix{Float64}:
 20.0  0.0  0.0
  1.0  9.0  0.0
  0.0  6.0 14.0

```

0.0.3 15.2

```

[1]: using LinearAlgebra
      using Plots
      include("lsq_classifier_data.jl");

```

(a)

```

[2]: function error_rate(X, y, theta)
      sz = size(X)[2]
      one = ones(sz)
      A = [X' one]
      y_hat = sign.(A*theta)
      correct = sum(y_hat.==y)
      return 1-correct/sz
    end

sz = size(X)[2];
one = ones(sz);
A = [X' one];
theta = pinv(A)*y;

```

```

er_train = error_rate(X, y, theta);
er_test = error_rate(X_test, y_test, theta);

println("Classification error on training: ", er_train, "\nClassification error_
↪on training: ", er_test)

```

Classification error on training: 0.19999999999999996
Classification error on training: 0.28

(b)

```

[3]: function regu(X, y, lambda)
    A = [X' ones(size(X)[2])]
    I = Array{Diagonal{ones(size(X)[1])}}
    newX = [A;sqrt(lambda)*I zeros(size(X)[1])]
    newY = [y;zeros(size(X)[1])]
    theta = newX\newy
    return theta
end

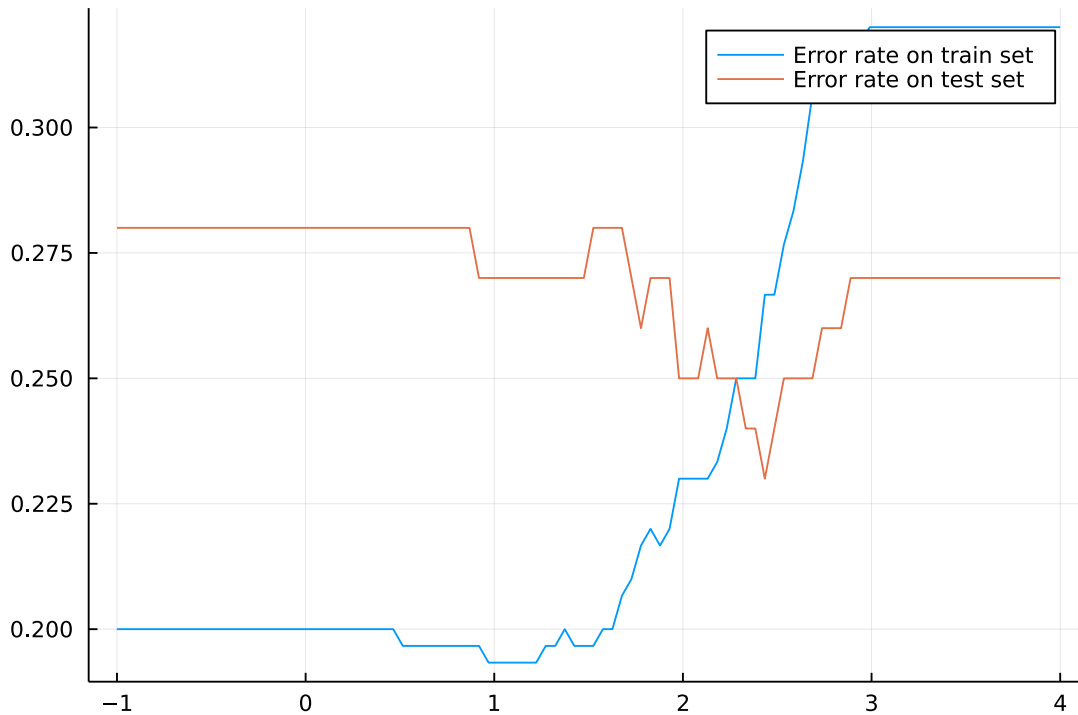
lambda = [10^i for i in range(-1,stop=4,length=100)];
xaxis = log10.(lambda);
error_train = zeros(100);
error_test = zeros(100);
for i in 1:100
    theta = regu(X, y, lambda[i]);
    error_train[i] = error_rate(X, y, theta);
    error_test[i] = error_rate(X_test, y_test, theta);
end
index = findmin(error_test)[2];
println("Reasonable lambda = ", lambda[index]);

plot(xaxis,error_train, label ="Error rate on train set")
plot!(xaxis,error_test, label ="Error rate on test set")

```

Reasonable lambda = 271.85882427329403

[3]:



0.0.4 15.3

```
[1]: using LinearAlgebra
include("price_elasticity.jl")
delta_p = (Prices .- p_nom) ./ p_nom;
delta_d = (Demands .- d_nom) ./ d_nom;
delta_p_train = delta_p[:,1:50];
delta_d_train = delta_d[:,1:50];
delta_p_test = delta_p[:,51:N];
delta_d_test = delta_d[:,51:N];

# Regression without regularization
E_hat = delta_d_train*pinv(delta_p_train)
```

```
[1]: 5×5 Matrix{Float64}:
-0.962267  -0.0833953  -0.520435  -0.648362  -0.387259
 0.107803  -0.30387   -0.0194254 -0.454458  -0.69952
-0.245873   0.0569617  -0.546831  -0.11042  -0.140529
-0.0147948 -0.298046   0.298149  -1.37377   0.794056
 0.360962  -0.912985  -1.39774   0.68058  -1.1817
```

Using regularization will avoid over-fitting problem.

```
[2]: # Regression with regularization
E_hat_regu = inv(delta_p_train*delta_p_train' + 0.
↳001*Array(Diagonal(ones(size(delta_p_train)[1]))))delta_p_train*delta_d_train'
```

```
[2]: 5×5 Matrix{Float64}:
-0.957856  0.106472  -0.244633  -0.0163845  0.358508
-0.0826826 -0.300272  0.0564209 -0.29443   -0.907341
-0.517686  -0.0195579 -0.543665  0.295951  -1.38978
-0.645188  -0.451344  -0.110264  -1.36318   0.676573
-0.385549  -0.695741  -0.139906  0.788808  -1.17367
```

```
[3]: test = norm(E_hat*delta_p_test-delta_d_test)/sqrt(n*25)
test_regu = norm(E_hat_regu*delta_p_test-delta_d_test)/sqrt(n*25)
println("RMS without regularization: ", test, "\nRMS using regularization: ",
↳test_regu)
```

RMS without regularization: 0.22099654989476042

RMS using regularization: 0.22175995430670603